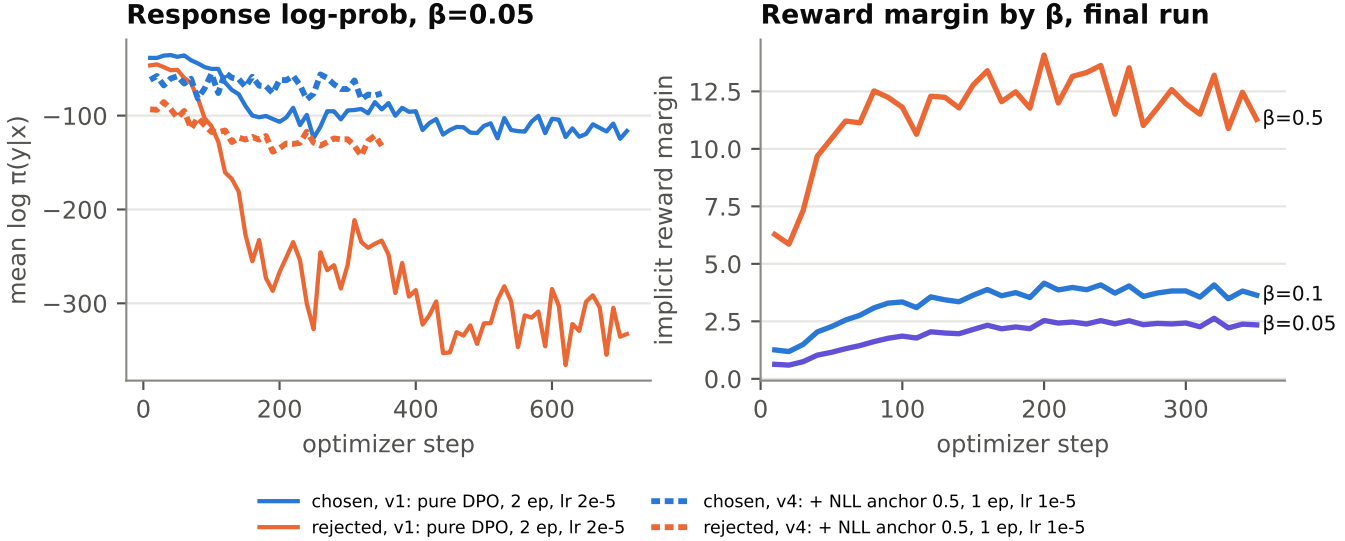
Response log-prob, β=0.05
mean log π(y|x)
−100
−200
−300
0
200
400
600
optimizer step
Reward margin by β, final run
implicit reward margin
12.5
10.0
7.5
5.0
2.5
0.0
0
100
200
300
optimizer step
β=0.5
β=0.1
β=0.05
chosen, v1: pure DPO, 2 ep, lr 2e-5
rejected, v1: pure DPO, 2 ep, lr 2e-5
chosen, v4: + NLL anchor 0.5, 1 ep, lr 1e-5
rejected, v4: + NLL anchor 0.5, 1 ep, lr 1e-5